

Johns Hopkins
Engineering for Professionals
605.767 Applied Computer Graphics

Brian Russin

Module 2B

Implementing a Ray Tracer




Ray Tracing: Intersections

- Primitive ray tracer checks intersection against every scene object
 - Looks for nearest intersection
 - Distance along the ray
 - Cannot simply stop search when an intersection has been found
 - Except for shadow rays
- Brute force strategy - recursive ray tracer spends most of its time testing for intersections
 - Each intersection spawns additional rays and subsequent intersect tests
 - Whitted estimated 95% in moderate complexity scenes
- For efficiency, need to address 2 problems:
 - Efficiently performing the intersection tests
 - Lecture 1
 - Strategy for guiding the order in which checks are performed
 - Limiting unnecessary intersection checks
 - Discuss some techniques at the end of this lecture



Ray-Tracing Spheres



Simple ray tracers often start with spheres placed in world coordinates. Simple intersection computation.

Intersecting Transformed Objects

- Wish to handle cases where an object is modeled in local coordinates and transformed to world coordinates
 - Translation, rotation, scaling (modeling transformation)
 - Transformation M maps an object to its image W
- Solution: solve for where the inverse transformed ray hits the original object
 - Rather than transform the object to world coordinates, transform the ray to the object's local coordinates
 - Less expensive!
 - Transform the ray origin by the inverse of the modeling transform
 - Transform the ray direction by the inverse of the modeling transform
 - Use a 0 as the homogenous factor – so the translate portion is ignored

$$r'(t) = M^{-1} \begin{bmatrix} R_x \\ R_y \\ R_z \\ 1 \end{bmatrix} + M^{-1} \begin{bmatrix} d_x \\ d_y \\ d_z \\ 0 \end{bmatrix}$$

Where R is the ray origin
and d is the ray direction




Handling Scaling

- Note that many ray intersection test methods assume a unit length direction vector for efficiency
 - For example: `Ray3::Intersect(BoundingSphere)`
- Any scaling within the modeling transformation **will impact the intersection distance (t)**
 - Need to “scale” t to counterbalance the change in distances the modeling transformation introduces
 - So we can recover the intersection point using the non-transformed ray



Transforming Normals

- Normal at the intersect point in local coordinates needs to be transformed to world coordinates
- Recall: normal vectors do not transform the same way as vertices
 - Non-uniform scaling in the modeling transform the normal vector will be distorted if transformed by the modeling transform



Scale by 2 in x direction, 1 in y

A 2D coordinate system with x and y axes. A purple square, horizontally stretched (width 2, height 1), is centered at the origin. A normal vector n is shown as an arrow pointing from the center towards the top-right corner. Below the diagram, the transformation of the normal vector is shown as a matrix multiplication:

$$n = \begin{bmatrix} 2 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 0.707 \\ 0.707 \\ 0 \\ 1 \end{bmatrix} = \begin{bmatrix} 1.414 \\ 0.707 \\ 0 \\ 1 \end{bmatrix}$$


Not Perpendicular to Surface!

Transforming Normals

- Think of normals as lying in a plane perpendicular to the surface
 - If p is a homogeneous plane (a,b,c,d) and v is a homogeneous vertex (x,y,z,w) on the plane, then $pv = 0$
 - $ax+by+cz+dw = 0$
 - Equivalent: $pM^{-1}Mv = 0$
 - If M is a nonsingular (has an inverse) vertex transformation matrix:
 - pM^{-1} is the **image of the plane** under the vertex transformation M
 - Plane must be transformed by premultiplication with the inverse of the modeling transform
- Alternatively, if v and n are perpendicular vectors
 - $n^T v = 0$
 - Definition of perpendicular vectors
 - For any nonsingular transformation M : $n^T M^{-1} M v = 0$
 - The transformed normal vector is: $(M^{-1})^T n$
 - **Normal vectors are transformed by the inverse transpose of the matrix that transform vertices**



Constructive Solid Geometry




Cassini Probe

Using BRL-CAD – an Open Source Solid Modeling package

<http://brlcad.org/>

CSG Representation

- Object stored in a tree structure
 - Leaves contain simple, geometric primitives
 - Nodes store operators and transformations
 - Figure 2.13 in Watt provides another example
- Can produce geometrically complex objects
 - Figure 2.14 in Watt
- Simple editing
 - Increase diameter of a bolt hole: increase cylinder radius
 - Much more difficult for boundary representation




Rendering CSG Representations

- CSG trades rendering efficiency for interactivity
 - Easy to interact and change the object
 - Interactive design
 - Inefficient rendering
- CAD / CAM systems until recently were primarily wireframe
 - Special off-line solid rendering procedure
- Primary rendering strategies include:
 - CSG ray-tracing
 - <http://web.cse.ohio-state.edu/~parent.1/classes/681/Lectures/19.RayTracingCSG.pdf>
 - Conversion to space subdivision data type
 - Volume rendering
 - Conversion to polygon mesh
 - Rendering with a stencil buffer
 - Requires a polygon mesh surface representation of each primitive



CSG Ray-Tracing

- Watt: Section 4.3, F.S. Hill Section 12.13
- Expensive method, but reasonably simple
 - Cast a ray from each pixel in the view plane
 - Parallel projection simplifies since rays are parallel
 - For each ray - test intersection with each primitive
 - Sort intersections based on Z depth
 - Consider boolean operations between first 2 primitives along the ray to find the first intersect point
 - Determine limits for composite object based on boolean operation
 - Apply shading model based on simple reflection model



Operation	Surface Limits
Union	A,D
Intersection	B,C
Difference (O2-O1)	C,D

Tracing a Ray Through the View Plane

$$H = N \tan(\theta/2)$$

$$W = H * \text{aspect}$$

Camera extends from $-W$ to $+W$ in u and $-H$ to $+H$ in v
 θ is the field of view angle in v direction
Aspect ratio is width over height of the framebuffer

To find the ray through the pixel with row r and column c :

$$r(t) = \text{eye} + \text{ray}_{rc} t$$

$$\text{ray}_{rc} = -Nn + W\left(\frac{2c}{n\text{Cols}} - 1\right)u + H\left(\frac{2r}{n\text{Rows}} - 1\right)v$$


Eye is the eye/camera position in world coordinates

View axes are u, v, n

N is the near plane distance (distance from the eye to viewplane)

$n\text{Rows}$, $n\text{Cols}$ are the image dimensions in pixels


Extend CameraNode class from 605.667 to support construction of a ray through a pixel




Drawing Pixel Blocks

- Ray casting/tracing can be time consuming
 - Would like to view scene prior to completion
 - Could render the scene with OpenGL
 - To verify camera and object positioning
- Sample lab renders “blocks” of pixels to help debug and test
 - Trace rays through one pixel, apply to a group of pixels
 - Builds a coarse approximation on first iteration
 - Successive iterations refine the image by reducing pixel block size
 - Can build a “memory framebuffer” to support pixel block drawing
 - Maintain a list of pixels that have been set
 - To avoid recalculating color at these pixels on successive iterations


16



4



1



Pseudo Code for Pixel Block Drawing

```
void display()
{
    MyFrameBuffer->Clear();           //Clear the memory framebuffer
    for (int blockSize = 16; blockSize > 0; blockSize /= 2)
    {
        for (int x = 0; x < imageWidth; x += blockSize)
        {
            for (int y = 0; y < imageHeight; y += blockSize)
            {
                if (!MyFrameBuffer->Set(x, y)) // Check if framebuffer has been set for this pixel
                {
                    // Construct a ray through the specified pixel and find the color
                    // using the ray caster or recursive ray tracer
                    MyFrameBuffer->Set(x, y, color, blockSize);
                }
            }
        }
        MyFrameBuffer->Render(); // Render the framebuffer
    }
}
```

