

Johns Hopkins
Engineering for Professionals
605.767 Applied Computer Graphics

Brian Russin

Module 8E

Shadow Maps



Shadow Z-Buffer or Shadow Mapping

- Lance Williams (1978) developed a shadow generation algorithm based on 2 passes through a Z-Buffer type algorithm
 - Lance Williams, “Casting Curved Shadows on Curved Surfaces,” SIGGRAPH 78
 - Image precision algorithm (operates in image space)
 - No knowledge of scene’s geometry is required
 - Pass1: for the light source(s)
 - Pass2: from the viewpoint
- A common **software** rendering technique
 - Pixar’s RenderMan uses the algorithm
 - Basic shadowing technique for Toy Story, etc.
- Hardware shadow mapping is now widely available
 - Uses texture mapping extensions
- Watt and Watt, Advanced Animation and Rendering Techniques, describes shadow Z-Buffer theory in extensive detail
 - Including pseudo-code and full anti-aliasing discussion
 - Haines and Moller discuss in section 7.4 (9.1.4 in 3rd Edition)



Shadow Z-Buffer: Creating the Shadow Map

- Shadow Z-Buffer is essentially depth testing from the light's point of view
 - Requires 2 passes
- Pass 1 produces the **shadow map** or **shadow Z-Buffer**
 - Render scene to store depth in the shadow Z-Buffer
 - 2D map indicating the depth of the closest pixels to the light
 - framebuffer object that can be used later as a texture
 - Uses the light source as the view point
 - This computes a depth from the light source
 - Disable writing to framebuffer
 - No lighting or texture mapping
 - The shadow map is used in the second pass
- Directional lights – orthographic projection
- Positional lights – perspective projection



Shadow Z-Buffer: Shadow Determination

- Pass 2: render the scene from the viewer viewpoint using a Z-Buffer algorithm with a couple enhancements:
 - If a pixel is visible (from regular Z-Buffer), it is transformed from screen space to the light source 'screen' space (x' , y' , z')
- Back map into world coordinates and project (including view transformation) using light source as the viewpoint
 - x' , y' provide an index into the shadow Z-Buffer
 - z' is compared to the value in the shadow Z-Buffer
 - If z' is greater, then another surface is nearer to the light source, point is in the shadow and a shadow intensity is assigned
 - Preferably the ambient intensity of the object being rendered
 - Otherwise, the point is on a surface nearest to the light source and is rendered normally
 - If z is approximately equal to z' we may have a case of “self-shadowing”
 - Object we are drawing is the one in the shadow Z-Buffer
 - Need to consider this case carefully




Shadow Z-Buffer (cont.)

- Shadow Z-Buffers are reasonably simple for single light source
 - Requires a separate shadow Z-Buffers for each light source
 - Or a single shadow Z-Buffers with a separate pass for each light source
 - Would be somewhat inefficient and slow
- Some drawbacks include:
 - Additional memory requirements
 - Per pixel transformation from screen space to light source 'screen' space
 - Shadow calculations performed for a pixel may be subsequently overwritten
- Shadow Z-Buffer is susceptible to aliasing
 - Similar to texture aliasing
 - Effectively projecting a pixel extent onto the shadow Z-Buffers
 - Many shadow map pixels map to one screen pixel
 - Pixel may be partly in shadow and partly not-yet we are making a binary decision
 - Anti-aliasing can be performed at substantial cost
 - Williams used filtering and dithering to reduce aliasing effects





Shadow Mapping

Object in shadow case ($A < B$)




Shadow Map Example




- Note the shadowing onto curved surfaces
 - Best technique for shadows onto curved surfaces
 - Notice how specular highlights never appear in shadows
- From Mark Kilgard presentation: “Shadow Mapping with Today’s OpenGL Hardware”.

Self Shadowing

- Regions where A approximately equal B are regions where self-shadowing is likely occurring
 - Fragment being drawn is the same as the closest to the light source
 - Round off in inverting projection/view as well as precision of buffers can lead to inequality
 - Want a slight “bias” when checking depth to prevent self-shadowing
 - Also known as “shadow acne”



Too little bias: everything self shadows



Too much bias: shadows can start too far away

Hardware Shadow Mapping Support

- Most hardware supports shadow mapping
 - Uses texture mapping resources
 - To store depth map
 - Perform filtering (for anti-aliasing)
 - Rely on projective texture
- NVidia GeForce, SGI, Xbox, and others
 - Performs the shadow test as a texture filtering operation
 - Modulate color with result
 - zero if fragment is shadowed or unchanged color if not

- OpenGL

```
glGenFramebuffers(1, &fbo);  
glBindFramebuffer(GL_FRAMEBUFFER, fbo);  
glGenTextures(1, &depth_tex);  
glBindTexture(GL_TEXTURE_2D, depth_tex);  
glTexImage2D(GL_TEXTURE_2D, 0, GL_DEPTH_COMPONENT16, w, h, 0, GL_DEPTH_COMPONENT, GL_FLOAT, 0);  
glFramebufferTexture(GL_FRAMEBUFFER, GL_DEPTH_ATTACHMENT, depth_tex, 0);
```



Pass 1: Constructing the Depth Map

- First pass: render scene from point of view of the light
 - Construct the framebuffer object and association depth texture
 - Performed once
 - Framebuffer object is reused every frame or when scene objects have been updated
 - Enable depth testing to find closest object to the light
 - Use **glPolygonOffset** to create bias in depth from light
 - Help prevent self shadowing
 - Usually better to error on the side of too much bias
 - e.g., `glPolygonOffset(factor = 1.1, bias = 4.0)`
 - or render only back-facing polygons
 - Adjust to suit the shadows in your scene
 - Depends somewhat on shadow map precision
 - Setup transformation from world space to light source space
 - Bind the framebuffer object and render the scene



Pass 2: Render Scene and Access the Depth Texture

- Render scene from eye's point-of-view as follows
 - Fragment's light position can be generated using the inverse View-Projection matrix multiplied by the Light projection matrix.
 - Set shadow map texture as uniform sampler
 - Set light projection as uniform variable
 - Multiply fragment position to get x' , y' , z' , w'
 - Sample the shadow map using x' , y' to get the shadow depth
 - Compare z' to the shadow depth value
 - Compute lighting based on whether the fragment is in shadow, or not
- Tutorial for OpenGL
 - <https://learnopengl.com/Advanced-Lighting/Shadows/Shadow-Mapping>



Percentage Closer Filtering

- Percentage Closer Filtering - extends shadow maps to provide pseudo soft shadows
 - Use texture filtering (like when texture magnification occurs) to determine whether the 4 nearest texels are in shadow or not
 - Bilinear interpolation used to estimate the how much light actually at the shadow location (pixel)
 - <http://graphics.pixar.com/library/ShadowMaps/paper.pdf>
 - http://developer.download.nvidia.com/shaderlibrary/docs/shadow_PCSS.pdf
- Other filtering methods can be used
 - Variance Shadow Maps
 - <http://developer.download.nvidia.com/SDK/10/direct3d/Source/VarianceShadowMapping/Doc/VarianceShadowMapping.pdf>



Shadow Maps with GLSL

- GLSL has shadow map sampling support
 - http://www.opengl.org/wiki/GLSL_Sampler
 - Uses **depth textures**
 - **shadow2D**, using a **sampler2DShadow** texture sampler
 - **shadow2DProj** – divides by coord.q
- Several references show how to do simple shadow maps with GLSL
 - <http://fabiansanglard.net/shadowmapping/index.php>



Issues with Shadow Mapping

- Prone to aliasing artifacts
 - “Percentage closer” filtering helps this
 - Normal color filtering does not work well
- Depth bias is not completely foolproof
 - Other methods like “second-depth shadow mapping” can work
 - Render only backfaces into shadow map
- Requires extra shadow map rendering pass and texture loading
- Higher resolution shadow map reduces blockiness
 - Increases texture copying expense
- Shadows are limited to view frustums
 - Not “all-directional”
 - Problematic for light sources in middle of scenes
- Objects outside or crossing the near and far clip planes are not properly accounted for by shadowing
 - Move near plane in as close as possible
 - Not so close that too much depth map precision is lost

