

Johns Hopkins
Engineering for Professionals
605.767 Applied Computer Graphics

Brian Russin

Module 11B

Locating Bottlenecks



Pipeline Optimization

- Primarily from From Haines and Moller
 - Chapter 18 (15 in 3rd Edition)
 - Highly recommended!
- Nvidia presentations
 - <https://docs.nvidia.com>
- In a pipeline architecture good performance depends on finding and eliminating **bottlenecks**
 - Bottleneck varies over the course of a frame
 - Find where the bottleneck resides over most of the frame



Performance Bottlenecks

- CPU / bus limited
 - Vertex and data transfer from CPU to graphics board is the limiting factor
 - Often related to how data is stored for use
 - Inefficient access
 - Application cannot feed data to the geometry stage fast enough
- Geometry limited
 - Performance is limited by large amount of geometry transforms being performed
 - Primitives too high resolution
 - Complex lighting
- Fill / raster limited
 - Large polygons with numerous pixel level operations
 - e.g. Blending can reduce performance



Rendering Speed

- Slowest pipeline stage determines the **rendering speed**
 - Expressed in frames per second
 - Number of rendered images per second
 - Slowest stage is called the **bottleneck**
- Framebuffer rates with double buffering
 - Rendering is synchronized with the update rate of the monitor
 - E.g. monitor with 72 Hz refresh will take $1/72 \text{ sec} = 0.01389 \text{ sec}$ to scan and display the framebuffer
 - Framebuffer can only be swapped during vertical retrace
 - Time between scan of bottom line of framebuffer and rewinding to top
 - If drawing takes longer than scan converter update rate it has to wait until next vertical retrace
 - Limits frame rates to scan rate / n
 - Example: 72Hz scan rate
 - Frame rates can be 72Hz, $72/2=36\text{Hz}$, $72/3=24\text{Hz}$, $72/4=18\text{Hz}$ etc.



Rendering Speed (cont.)

- Use single buffering to measure total rendering time
- Average frame rate is unreliable measurement
 - e.g. Three seconds are analyzed with 50fps, 50fps, 20fps.
 - Average is **40fps**
- **Incorrect** when measuring render time
 - 50fps = 20 ms / frame
 - 20fps = 50 ms / frame
 - Average: $(20 + 20 + 50)/3 = 30$ ms / frame = **33.3 fps!**



Application Stage

- Developer controls the application stage
 - Executes in software
- Two main phases
 - Main process and subsystems
 - Update scene graph elements
 - Culling: view frustum, occlusion, contribution
 - Collision detection
 - Animation: moving objects, texture animation, morphing
 - Graphics API calls
 - Feeds geometry to the geometry stage
 - Points, lines, triangles
 - Data access, processing/culling, etc.
 - Updates to uniform buffers, textures, etc.



Pipeline Optimization

- 3 stage process:
 - 1) Locate the bottleneck
 - 2) Perform some optimization on the bottleneck stage
 - Decreases the workload in the bottleneck stage
 - 3) Repeat step 1 if performance goal not achieved
 - Note: location of bottleneck may have changed!
- Do not spend too much effort on 2 (over-optimization)
 - If bottleneck moves you may waste effort
- When bottleneck stage cannot be further optimized you can increase work in other stages
 - Will not decrease frame rate since bottleneck stage is not altered
 - Examples
 - Use more complex lighting if geometry stage is not bottleneck
 - Add textures or blending if rasterization stage is not bottleneck



Locating the Bottleneck

- Cannot find bottleneck by simply timing the process
 - Tells total time for bottleneck stage, not which stage
- Must set up specific tests that affect one stage at a time
 - When total rendering time is lowered we have found the bottleneck
- Related method: reduce workload in other stages
 - Do not reduce workload on the stage being tested
 - If performance does not change then the bottleneck is in the stage not altered



Testing the Application Stage

- Use a utility to measure CPU work
 - UNIX: **top** or **osview**
 - Windows Task Manager
 - If near 100% CPU usage the application is likely the bottleneck
 - Caution: could be a **busy wait** condition
- Use a **null driver**
 - Accepts API calls but does nothing
 - Geometry and rasterization stages will do little or no work
 - Provides an upper limit on how fast the application can run
 - Can hide bottleneck that is due to communication between application and geometry stage



Testing the Geometry Stage

- Most difficult stage to test
 - Change of workload here can change load in other stages as well
- Lighting operations only affect the geometry stage
 - Disable or add light sources
 - Test is performance is impacted
- Increase/Decrease the vertex buffer sizes
- Add logic to vertex shader
 - Although changes that do nothing may be optimized away
- Eliminating the other stages
 - Test application and rasterization stages
 - If neither is the bottleneck then the geometry stage must be!



Testing the Rasterization Stage

- Easy test
- Decrease the resolution of the window/ viewport
 - Fewer pixels are touched – less texture / depth buffer tests
 - If rendering performance increases then the pipeline is **fill-limited**
 - Rasterizer stage is the bottleneck
- Other tests:
 - Turn off depth testing and blending
 - If performance increases then rasterization stage is the bottleneck
 - Vary texture sizes or filtering options
 - May be “fill-limited” due to texture access/filtering

